

Homework #3

CSE 446: Machine Learning

Eric Boris: 1976637

Conceptual Questions

Problem 1

- a. **False:** SVM only maximizes the margin, it doesn't minimize generalization error among linear classifiers.
- b. **Decrease:** Lower values of σ allow the model to be more expressive.
- c. **True:** Neural networks often have non-convex loss functions with only local minima.
- d. **False:** This can lead to problems. It's better to initialize weights to random values.
- e. **True:** Without nonlinearities in the network, the network can only learn linear functions.
- f. **True:** We mitigate this by using SGD to improve the runtime of the backward pass.

Kernels

Problem 2

Show that $K(x, x') = e^{-\frac{(x-x')^2}{2}}$ is a kernel function for this feature map, i.e., $\phi(x) \cdot \phi(x') = e^{-\frac{(x-x')^2}{2}}$.

$$\begin{aligned}\phi(x) \cdot \phi(x') &= \left(\sum_{i=0}^{\infty} \frac{1}{\sqrt{i!}} e^{-\frac{x^2}{2}} x^i \right) \left(\sum_{i=0}^{\infty} \frac{1}{\sqrt{i!}} e^{-\frac{x'^2}{2}} x'^i \right) \\ &= e^{-\frac{x^2+x'^2}{2}} \sum_{i=0}^{\infty} \left(\frac{1}{i!} (xx')^i \right) \\ &= e^{-\frac{x^2+x'^2}{2}} e^{xx'} && \text{Taylor Series of } e^x = \sum_{i=0}^{\infty} \frac{x^i}{i!} \\ &= e^{-\frac{(x-x')^2}{2}}\end{aligned}$$

Problem 3: Answers

- a. Polynomial Kernel: Lambda=1E-05 d=11.0
RBF Kernel: Lambda=1E-03 gamma=2.0802
- b. See Figures 1 and 2

Problem 3: Code

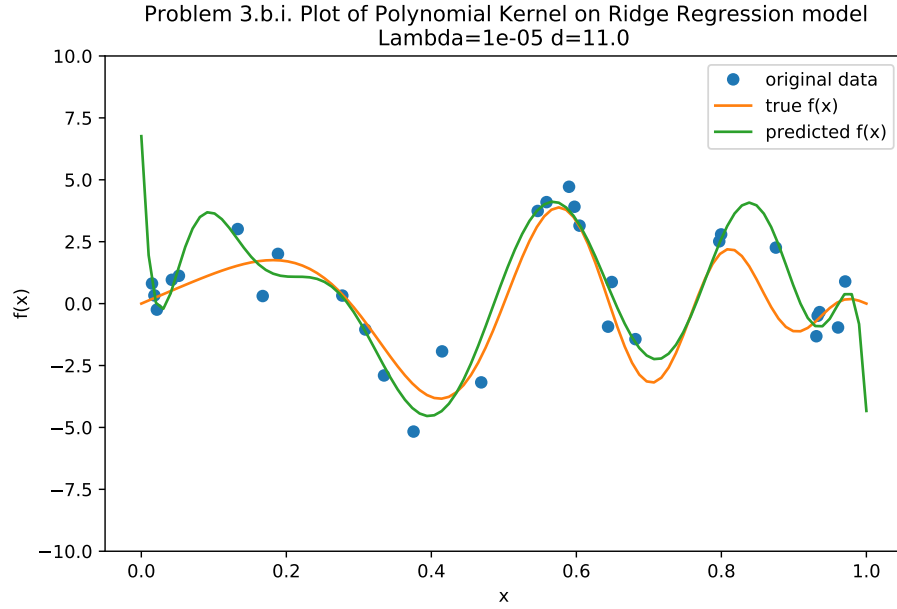


Figure 1




Figure 2

```

1 # Problem 3: Kernel Ridge Regression
2
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 def generate_data(f, n):
7     ''' Generate n random samples of data and actual output using function f. '''
8     # Let  $x_i$  be uniformly random on  $[0, 1]$ .
9     x = np.random.rand(n)
10    # Let  $\epsilon_i \sim N(0, 1)$ .

```

```

11     epsilon = np.random.randn(n)
12     # Let  $y_i = f(x_i) + \epsilon_i$ .
13     y = f(x) + epsilon
14     # Return x as a column vector.
15     return x.reshape(-1, 1), y
16
17 def f_star(x):
18     ''' Compute the f star function given in the spec. '''
19     return 4 * np.sin(np.pi * x) * np.cos(6 * np.pi * x ** 2)
20
21 def polynomial_kernel(x, z, d):
22     ''' Define the polynomial kernel given in the spec where d is a hyperparameter. '''
23     return (1 + x.dot(z.T)) ** d
24
25 def rbf_kernel(x, z, gamma):
26     ''' Define the RBF kernel given in the spec where gamma is a hyperparameter. '''
27     return np.exp(-gamma * squared_difference(x, z))
28
29 def squared_difference(x, z):
30     ''' Return the squared difference between x and z. '''
31     return np.sum((x[:, :, None] - z[:, :, None].T) ** 2, axis=1)
32
33 def leave_one_out_cv(model, X, y, regularized_lambdas, hyperparameters):
34     ''' Compute the error of lambda and hypermeter combinations on the model using LOOCV. '''
35
36     def error(model, X, y, regularized_lambda, hyperparameter):
37         ''' Perform LOOCV on one lambda / hyperparameter pair and return the model mean error.
38             '''
39         # Set the model to use the parameters
40         model.regularized_lambda = regularized_lambda
41         model.hyperparameter = hyperparameter
42
43         # Perform LOOCV.
44         error = []
45         for i in range(len(X)):
46             # Use this to determine which indices to include in the "in" subsets.
47             indices = np.full(len(X), True)
48             indices[i] = False
49
50             # Define the cross validation subsets.
51             X_in, y_in = X[indices], y[indices]
52             X_out, y_out = X[i].reshape(1, -1), y[i].reshape(1, )
53
54             # Interpolate with the model.
55             model.train(X_in, y_in)
56             y_hat = model.predict(X_out)
57
58             # Compute and store the mean squared error.
59             error.append(np.mean((y_hat - y_out) ** 2))
60
61         return np.mean(error)
62
63     # Perform LOOCV.
64     results = [(error(model, X, y, rl, hp), rl, hp) for rl in regularized_lambdas for hp in
65                 hyperparameters]
66     return np.array(results)
67
68 def plot(title, subtitle, x_label, y_label, file_path, original_x, original_y, original_label,
69         true_x, true_y, true_label, pred_y, pred_label, argsort, dim=(8, 5)):
70     ''' Plot the graphs for Problem 3 Part b. '''
71     plt.figure(figsize=dim)
72     plt.title(f'{title}\n{subtitle}')
73     plt.xlabel(x_label)
74     plt.ylabel(y_label)
75     plt.plot(original_x[argsort, 0], original_y[argsort], 'o', label=original_label)
76     plt.plot(true_x[:, 0], true_y, label=true_label)
77     plt.plot(true_x[:, 0], pred_y, label=pred_label)
78     plt.legend()
79     plt.ylim([-10, 10])

```

```

77     plt.savefig(file_path)
78     plt.show()
79
80 class KernelRidgeRegression:
81     def __init__(self, kernel, regularized_lambda=1E-8, hyperparameter=None):
82         self.kernel = kernel
83         self.regularized_lambda = regularized_lambda
84         self.hyperparameter = hyperparameter
85
86     def train(self, X_train, y_train):
87         ''' Train the model. '''
88         self.mean = np.mean(X_train, axis=0)
89         self.std = np.std(X_train, axis=0)
90
91         X_train = (X_train - self.mean) / self.std
92         self.X_train = X_train
93
94         K = self.kernel(X_train, X_train, self.hyperparameter)
95         self.alpha = np.linalg.solve(K + self.regularized_lambda * np.eye(K.shape[0]), y_train)
96
97         return self.alpha
98
99     def predict(self, X):
100         ''' Predict using the model. '''
101         X = (X - self.mean) / self.std
102         K = self.kernel(self.X_train, X, self.hyperparameter)
103         return K.T.dot(self.alpha)
104
105 def main():
106     # Generate training data.
107     X_train, y_train = generate_data(f=f_star, n=30)
108
109     # Set the hyperparameter bounds.
110     regularized_lambdas = 10.0 ** (-np.arange(2, 10))
111     ds = np.arange(4, 20)
112     gammas = (1 / np.median(squared_difference(X_train, X_train))) * np.linspace(0, 2, 10)
113
114     # Build and score both models.
115     poly_model = KernelRidgeRegression(polynomial_kernel)
116     poly_results = leave_one_out_cv(poly_model, X_train, y_train, regularized_lambdas, ds)
117
118     rbf_model = KernelRidgeRegression(rbf_kernel)
119     rbf_results = leave_one_out_cv(rbf_model, X_train, y_train, regularized_lambdas, gammas)
120
121     # Part a. Find and assign the best lambdas and hyperparameters for and to the models.
122     poly_model.regularized_lambda = poly_results[np.argmax(poly_results[:, 0])][1]
123     poly_model.hyperparameter = poly_results[np.argmax(poly_results[:, 0])][2]
124
125     rbf_model.regularized_lambda = rbf_results[np.argmax(rbf_results[:, 0])][1]
126     rbf_model.hyperparameter = rbf_results[np.argmax(rbf_results[:, 0])][2]
127
128     print(f'Problem 3.a.i. Best Lambda and d values with polynomial kernel: Lambda={poly_model.regularized_lambda}\td={poly_model.hyperparameter}')
129     print(f'Problem 3.a.ii. Best Lambda and gamma values with RBF kernel: Lambda={rbf_model.regularized_lambda}\tgamma={rbf_model.hyperparameter}')
130
131     # Part b. Plot the learned functions using the best lambdas and hyperparameters from part a.
132     X_evenly_spaced = np.linspace(0, 1, 100).reshape(-1, 1)
133     y_evenly_spaced = f_star(X_evenly_spaced[:, 0])
134     argsort = np.argsort(X_train[:, 0])
135
136     poly_model.train(X_train, y_train)
137     y_hat_evenly_spaced_poly = poly_model.predict(X_evenly_spaced)
138
139     plot(title='Problem 3.b.i. Plot of Polynomial Kernel on Ridge Regression model',
140         subtitle=f'Lambda={poly_model.regularized_lambda} d={poly_model.hyperparameter}',
141         xlabel='x',

```

```

142     y_label='f(x)',
143     file_path='../plots/3bi.pdf',
144     original_x=X_train,
145     original_y=y_train,
146     original_label='original data',
147     true_x=X_evenly_spaced,
148     true_y=y_evenly_spaced,
149     true_label='true f(x)',
150     pred_y=y_hat_evenly_spaced_poly,
151     pred_label='predicted f(x)',
152     argsort=argsort)
153
154 rbf_model.train(X_train, y_train)
155 y_hat_evenly_spaced_rbf = rbf_model.predict(X_evenly_spaced)
156
157 plot(title='Problem 3.b.ii. Plot of RBF Kernel on Ridge Regression model',
158      subtitle=f'Lambda={rbf_model.regularized_lambda} gamma={rbf_model.hyperparameter}',
159      x_label='x',
160      y_label='f(x)',
161      file_path='../plots/3bii.pdf',
162      original_x=X_train,
163      original_y=y_train,
164      original_label='original data',
165      true_x=X_evenly_spaced,
166      true_y=y_evenly_spaced,
167      true_label='true f(x)',
168      pred_y=y_hat_evenly_spaced_rbf,
169      pred_label='predicted f(x)',
170      argsort=argsort)
171
172 if __name__ == '__main__':
173     main()

```

Neural Networks for MNIST

Problem 4: Answers

a., b., c.

We see that a deep network performs significantly better than a wide network with the same number of parameters. The deep network reaches comparable levels of performance to the wide in one epoch compared to eighteen epochs. This performance is shown in the following tables and plotted in figures 3 and 4. The reason for the deep network's better performance is likely due to there being more nonlinearities in the deeper model so the data are more immediately fit.

Model Performance

	Wide Net		Deep Net	
	Training	Testing	Training	Testing
Accuracy	0.9904	0.9749	0.9911	0.9718
Loss	0.0336	0.0844	0.0320	0.0928

Model Details

	Wide Net	Deep Net
Number of Parameters	50890	50890
Epochs trained	18	1

Problem 4: Code

```

1 # Problem 4: Neural Networks for MNIST
2
3 import numpy as np

```

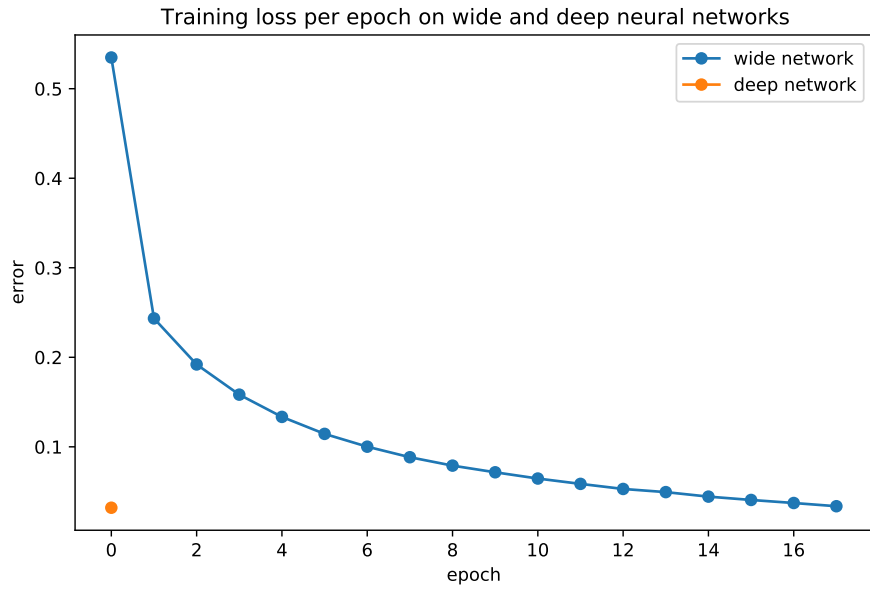


Figure 3

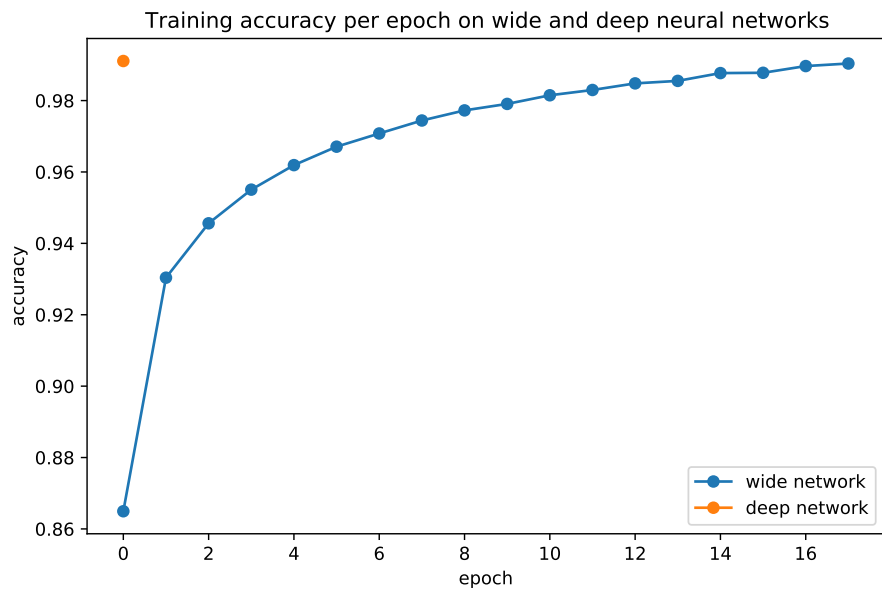


Figure 4

```

4 import torch
5 import torch.nn as nn
6 import torchvision.datasets as datasets
7 import torchvision.transforms as transforms
8 import matplotlib.pyplot as plt
9 from tqdm import tqdm
10
11 def get_loaders(root_path):
12     ''' Return MNIST train and test DataLoaders. '''
13     train = torch.utils.data.DataLoader(

```

```

14         datasets.MNIST(
15             root=root_path,
16             train=True,
17             download=True,
18             transform=transforms.ToTensor()),
19         batch_size=128,
20         shuffle=True)
21     test = torch.utils.data.DataLoader(
22         datasets.MNIST(
23             root=root_path,
24             train=False,
25             download=True,
26             transform=transforms.ToTensor()),
27         batch_size=128,
28         shuffle=True)
29     return train, test
30
31 def count_parameters(weights, biases):
32     ''' Return the total number of parameters used in the given weights and biases. '''
33     parameters = sum([np.prod(w.shape) for w in weights])
34     parameters += sum([np.prod(b.shape) for b in biases])
35     return parameters
36
37 def plot(title, x_label, y_label, file_name, x_1, y_1, label_1, x_2, y_2, label_2,
38         dim=(8, 5)):
39     ''' Plot two lines on single chart. '''
40     plt.figure(figsize=dim)
41     plt.title(title)
42     plt.xlabel(x_label)
43     plt.ylabel(y_label)
44     plt.plot(x_1, y_1, '-o', label=label_1)
45     plt.plot(x_2, y_2, '-o', label=label_2)
46     plt.xticks(np.arange(0, max(x_1)+1, 2.0))
47     plt.legend()
48     plt.savefig(file_name)
49     plt.show()
50
51 class NeuralNetwork:
52     def __init__(self, data_loader, learning_rate, n_neurons, n_layers, input_dim,
53                 output_dim):
54         self.data_loader = data_loader
55         self.learning_rate = learning_rate
56
57         alpha = 1 / (np.sqrt(input_dim))
58         self._weights(alpha, n_neurons, n_layers, input_dim, output_dim)
59         self._biases(alpha, n_neurons, n_layers, output_dim)
60
61     def _weights(self, alpha, n_neurons, n_layers, input_dim, output_dim):
62         ''' Create the model's weights. '''
63         weights = []
64
65         # Define the input layer weights.
66         weights.append(-2 * alpha * torch.rand(n_neurons, input_dim) + alpha)
67         weights[-1].requires_grad = True
68
69         # Define the hidden layer weights.
70         # Don't add the input and output layer weights.
71         for _ in range(n_layers - 2):
72             weights.append(-2 * alpha * torch.rand(n_neurons, n_neurons) + alpha)
73             weights[-1].requires_grad = True
74
75         # Define the output layer weights.
76         weights.append(-2 * alpha * torch.rand(output_dim, n_neurons) + alpha)
77         weights[-1].requires_grad = True
78
79         self.weights = weights
80
81     def _biases(self, alpha, n_neurons, n_layers, output_dim):

```

```

80     ''' Create the model's biases. '''
81     biases = []
82
83     # Define the input and hidden layer biases.
84     # Don't add the output layer biases.
85     for _ in range(n_layers - 1):
86         biases.append(-2 * alpha * torch.rand(n_neurons) + alpha)
87         biases[-1].requires_grad = True
88
89     # Define the output layer biases.
90     biases.append(-2 * alpha * torch.rand(output_dim) + alpha)
91     biases[-1].requires_grad = True
92
93     self.biases = biases
94
95     def train(self, n_epochs, accuracy_threshold=0.99, verbose=False):
96         ''' Train the model. '''
97         # Use the weights and biases as the parameter.
98         optimizer = torch.optim.Adam(self.weights + self.biases, lr=self.
99             learning_rate)
100
101         accuracies = []
102         losses = []
103         for epoch in range(n_epochs):
104             # Perform forward and backwards passes through the data and return the
105             model performance.
106             accuracy, loss = self.measure_performance(self.data_loader, optimizer=
107                 optimizer, train=True)
108
109             accuracies.append(accuracy)
110             losses.append(loss)
111
112             if verbose:
113                 print(f'Epoch={epoch}\tTraining Loss={losses[-1]}\tAccuracy={accuracy}
114                     ')
115
116             # End training if minimum accuracy threshold is met.
117             if accuracy > accuracy_threshold:
118                 break
119
120         return accuracies, losses, self.weights, self.biases
121
122     def measure_performance(self, data_loader, optimizer=None, train=False):
123         ''' Perform forwards and backwards passes on the model and return the
124         performance. '''
125         accuracy = 0
126         loss = 0
127
128         for X, y in tqdm(iter(data_loader)):
129             # Change the dimensions of X.
130             X = torch.flatten(X, start_dim=1, end_dim=3)
131
132             # Perform the forward pass and get the predictions.
133             logits = self._forward(X)
134             y_hat = torch.argmax(logits, 1)
135
136             # Compute the accuracy and loss.
137             accuracy += torch.sum(y == y_hat)
138             loss_tmp = torch.nn.functional.cross_entropy(logits, y, size_average=
139                 False)
140
141             # Gradient descent backward pass.
142             if train:
143                 optimizer.zero_grad()
144                 loss_tmp.backward()
145                 optimizer.step()
146
147             loss += loss_tmp

```



```

143         # Normalize the performance measures.
144         loss /= len(data_loader.dataset)
145         accuracy = accuracy.to(dtype=torch.float) / len(data_loader.dataset)
146
147         return accuracy, loss
148
149     def _forward(self, x):
150         ''' Perform a pass through the network using the given input x and ReLU
151             nonlinearities. '''
152         y = torch.matmul(x, self.weights[0].T) + self.biases[0]
153         for i in range(1, len(self.weights)):
154             y = torch.matmul(nn.functional.relu(y), self.weights[i].T) + self.biases[
155                 i]
156         return y
157
158     def main():
159         # Load the training and test data.
160         train_loader, test_loader = get_loaders('../data/python_mnist/')
161
162         # Part a: Build, train, and test a wide neural network.
163         wide_net = NeuralNetwork(data_loader=train_loader, learning_rate=1E-3, n_neurons
164                                 =64, n_layers=2, input_dim=784, output_dim=10)
165
166         wide_train_accuracies, wide_train_losses, wide_weights, wide_biases = wide_net.
167             train(n_epochs=500, verbose=True)
168         wide_parameters = count_parameters(wide_weights, wide_biases)
169         print(f'Wide net training results:\nAccuracy={wide_train_accuracies[-1]}\tLoss={
170             wide_train_losses[-1]}\tN Parameters={wide_parameters}\n')
171
172         wide_test_accuracy, wide_test_loss = wide_net.measure_performance(test_loader)
173         print(f'Wide net test results:\nAccuracy={wide_test_accuracy}\tLoss={
174             wide_test_loss}\n')
175
176         # Part b: Build, train, and test a deep neural network.
177         deep_net = NeuralNetwork(data_loader=train_loader, learning_rate=1E-3, n_neurons
178                                 =32, n_layers=3, input_dim=784, output_dim=10)
179
180         deep_train_accuracies, deep_train_losses, deep_weights, deep_biases = wide_net.
181             train(n_epochs=500, verbose=True)
182         deep_parameters = count_parameters(deep_weights, deep_biases)
183         print(f'Deep net training results:\nAccuracy={deep_train_accuracies[-1]}\tLoss={
184             deep_train_losses[-1]}\tN Parameters={deep_parameters}\n')
185
186         deep_test_accuracy, deep_test_loss = wide_net.measure_performance(test_loader)
187         print(f'Deep Net Test Results:\nAccuracy={deep_test_accuracy}\tLoss={
188             deep_test_loss}\n')
189
190         # Plot the performance of both models.
191         x_wide_evenly_spaced = range(len(wide_train_losses))
192         x_deep_evenly_spaced = range(len(deep_train_losses))
193
194         plot(title='Training loss per epoch on wide and deep neural networks',
195              x_label='epoch',
196              y_label='error',
197              file_name='../plots/4_losses.pdf',
198              x1=x_wide_evenly_spaced,
199              y1=[float(loss) for loss in wide_train_losses],
200              label_1='wide network',
201              x2=x_deep_evenly_spaced,
202              y2=[float(loss) for loss in deep_train_losses],
203              label_2='deep network')
204
205         plot(title='Training accuracy per epoch on wide and deep neural networks',
206              x_label='epoch',
207              y_label='accuracy',
208              file_name='../plots/4_accuracies.pdf',
209              x1=x_wide_evenly_spaced,
210              y1=wide_train_accuracies,
211              label_1='wide network',

```

```

202         x_2=x_deep_evenly_spaced ,
203         y_2=deep_train_accuracies ,
204         label_2='deep network')
205
206
207 if __name__ == '__main__':
208     main()

```

Using Pretrained Networks and Transfer Learning

Problem 5: Answers

a. Fixed Feature Extractor:

- Plot for training loss: See figure 5




Figure 5

- Plot for validation loss: See figure 6
- Highest validation accuracy: 0.8046
- Test accuracy: 79.99
- Test loss: 0.6293

b. Fine-Tuning:

- Plot for training loss: See figure 7
- Plot for validation loss: See figure 8
- Highest validation accuracy: 0.9088
- Test accuracy: 90.73




Figure 6




Figure 7

- Test loss: 0.4235

Problem 5: Code

```

1 # -*- coding: utf-8 -*-
2 """
3 CSE446 hw3 p5.ipynb

```

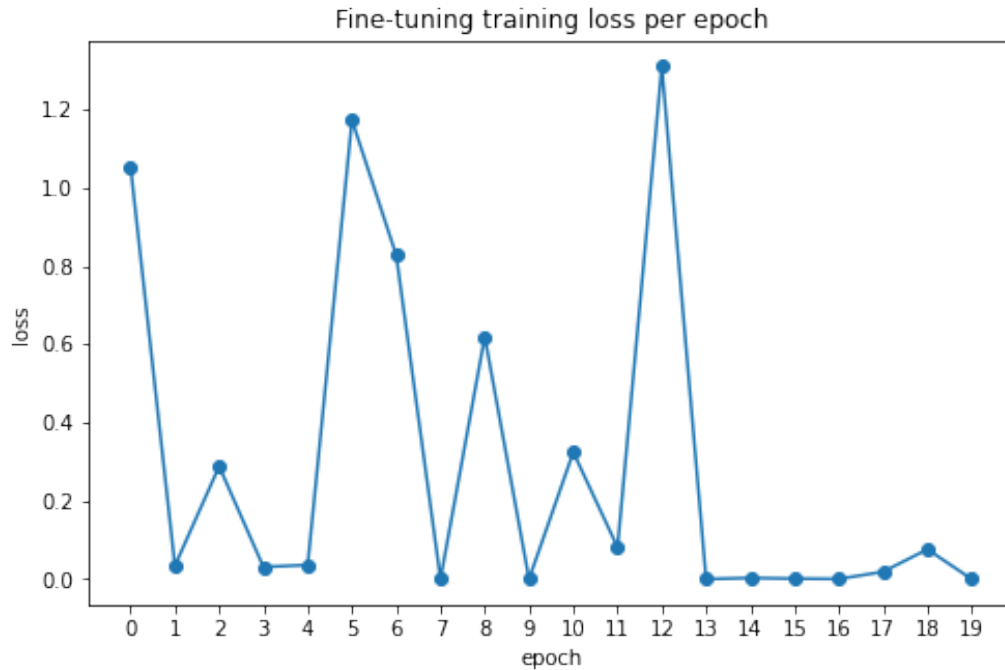


Figure 8

```

4
5 Automatically generated by Colaboratory.
6
7 Original file is located at
8 https://colab.research.google.com/drive/1EtJYFeZDQPBH3zOfZk9kJQ2KNAagt9FT
9
10 Written using code from the following tutorials:
11 https://pytorch.org/tutorials/beginner/blitz/cifar10\_tutorial.html
12 https://pytorch.org/tutorials/beginner/transfer\_learning\_tutorial.html#load-data
13 https://pytorch.org/tutorials/beginner/finetuning\_torchvision\_models\_tutorial.html
14 """
15
16 # Access google drive for saving and loading trained models.
17 from google.colab import drive
18 drive.mount('/content/drive')
19
20 import torch
21 import torchvision
22 import torchvision.transforms as transforms
23 import matplotlib.pyplot as plt
24 import numpy as np
25 import torch.nn as nn
26 import torch.optim as optim
27 from tqdm import tqdm
28 from torch.utils.data import random_split
29 import time
30 import copy
31
32 # Use GPU if it's available, and CPU otherwise.
33 train_on_gpu = torch.cuda.is_available()
34 train_on_multi_gpu = (torch.cuda.device_count() >= 2)
35 gpus = torch.cuda.device_count()
36
37 # Define transforms for the data to work with the AlexNet model.
38 transform = transforms.Compose([transforms.Resize(256),
39                                transforms.ToTensor(),
40                                transforms.Normalize((0.5, 0.5, 0.5),
41                                                    (0.5, 0.5, 0.5))])

```

```

42
43 # Load the dataset.
44 dataset = torchvision.datasets.CIFAR10(root='./data', train=True,
45                                       download=True, transform=transform)
46
47 # Split the dataset into training / validation such that validation is 10% of
48 # the training.
49
50 # Use a seed so that the training / validation split is the same each time.
51 torch.manual_seed(43)
52
53 # We set the validation set to be 10% of the training data.
54 val_size = 5000
55 train_size = len(dataset) - val_size
56
57 # Split the set.
58 trainset, valset = random_split(dataset, [train_size, val_size])
59
60 # Define the training and validation dataloaders.
61 trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
62                                           shuffle=True, num_workers=2)
63
64 valloader = torch.utils.data.DataLoader(valset, batch_size=4,
65                                         shuffle=True, num_workers=2)
66
67 # Join train and validation into a single dataloader
68 dataloaders = {'train': trainloader, 'val': valloader}
69 dataset_sizes = {'train': train_size, 'val': val_size}
70
71 # Load and define the test dataloader.
72 testset = torchvision.datasets.CIFAR10(root='./data', train=False,
73                                       download=True, transform=transform)
74 testloader = torch.utils.data.DataLoader(testset, batch_size=4,
75                                         shuffle=False, num_workers=2)
76 test_size = len(testset)
77
78 # Define the dataset class names.
79 classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
80           'ship', 'truck')
81
82 def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
83     ''' Train the model and return the trained model as well as training
84         performance data. '''
85
86     # Handle training on gpu or cpu.
87     if train_on_multi_gpus:
88         print(f"\nTraining on {gpus} GPUs!\n")
89         model = torch.nn.DataParallel(model).cuda()
90     elif train_on_gpu:
91         print('\nTraining on GPU!\n')
92         model = model.cuda()
93     else:
94         print('\nTraining on CPU; consider making n-epochs very small.\n')
95
96     since = time.time()
97
98     best_model_wts = copy.deepcopy(model.state_dict())
99     best_acc = 0.0
100
101     train_losses = []
102     valid_losses = []
103
104     for epoch in tqdm(range(num_epochs)):
105         # Each epoch has a training and validation phase
106         for phase in ['train', 'val']:
107             if phase == 'train':
108                 model.train() # Set model to training mode
109             else:
110                 model.eval() # Set model to evaluate mode

```

```

111
112     running_loss = 0.0
113     running_corrects = 0
114
115     # Iterate over data.
116     for inputs, labels in dataloaders[phase]:
117         if train_on_multi_gpu or train_on_gpu:
118             inputs, labels = inputs.cuda(), labels.cuda()
119
120         # zero the parameter gradients
121         optimizer.zero_grad()
122
123         # forward
124         # track history if only in train
125         with torch.set_grad_enabled(phase == 'train'):
126             outputs = model(inputs)
127             _, preds = torch.max(outputs, 1)
128             loss = criterion(outputs, labels)
129
130             # backward + optimize only if in training phase
131             if phase == 'train':
132                 loss.backward()
133                 optimizer.step()
134
135             # statistics
136             running_loss += loss.item() * inputs.size(0)
137             running_corrects += torch.sum(preds == labels.data)
138     if phase == 'train':
139         scheduler.step()
140
141     epoch_loss = running_loss / dataset_sizes[phase]
142     epoch_acc = running_corrects.double() / dataset_sizes[phase]
143
144     print('{} Loss: {:.4f} Acc: {:.4f}'.format(
145         phase, epoch_loss, epoch_acc))
146
147     # deep copy the model
148     if phase == 'val' and epoch_acc > best_acc:
149         best_acc = epoch_acc
150         best_model_wts = copy.deepcopy(model.state_dict())
151
152     # Store this epoch's training and validation losses.
153     if phase == 'train':
154         train_losses.append(loss)
155     else:
156         valid_losses.append(loss)
157
158     time_elapsed = time.time() - since
159     print('Training complete in {:.0f}m {:.0f}s'.format(
160         time_elapsed // 60, time_elapsed % 60))
161     print('Best val Acc: {:.4f}'.format(best_acc))
162
163     # load best model weights
164     model.load_state_dict(best_model_wts)
165     return model, train_losses, valid_losses
166
167 def performance(model):
168     ''' Return the loss and performace of the model on the test data. '''
169     correct = 0
170     total = 0
171     running_loss = 0.0
172     with torch.no_grad():
173         for data in testloader:
174             images, labels = data
175
176             # Handle for gpu.
177             if train_on_multi_gpu or train_on_gpu:
178                 images, labels = images.cuda(), labels.cuda()
179

```

```

180         # Compute the accuracy.
181         outputs = model(images)
182         _, predicted = torch.max(outputs.data, 1)
183         total += labels.size(0)
184         correct += (predicted == labels).sum().item()
185
186         # Compute the loss.
187         loss = criterion(outputs, labels)
188         running_loss += loss.item() * images.size(0)
189
190     loss = running_loss / test_size
191     accuracy = 100 * correct / total
192
193     return loss, accuracy
194
195 def plot(title, x_label, y_label, x, y, file_dir, dim=(8, 5)):
196     ''' Plot the given data. '''
197     plt.figure(figsize=dim)
198     plt.title(title)
199     plt.xlabel(x_label)
200     plt.xticks(np.arange(0, max(x)+2, 1))
201     plt.ylabel(y_label)
202     plt.plot(x, y, '-o')
203     plt.show()
204     plt.savefig(file_dir)
205
206 # Define the fixed feature extractor model and it's related components.
207 ffe_model = torchvision.models.alexnet(pretrained=True)
208
209 # Prevent all but the last layer from training.
210 for param in ffe_model.parameters():
211     param.requires_grad = False
212
213 ffe_model.classifier[6] = nn.Linear(4096, 10)
214
215 criterion = nn.CrossEntropyLoss()
216 optimizer = optim.SGD(ffe_model.parameters(), lr=0.001, momentum=0.9)
217 scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)
218
219 # Train and save the ffe model.
220 ffe_model, ffe_train_losses, ffe_valid_losses = train_model(ffe_model, criterion, optimizer,
221     scheduler, 20)
222
223 PATH = '/content/drive/MyDrive/CSE446_hw3/cifar_ffe_model.pth'
224 torch.save(ffe_model.state_dict(), PATH)
225
226 plot(title='Fixed feature extractor training loss per epoch',
227     x_label='epoch',
228     y_label='loss',
229     x=[x for x in range(len(ffe_train_losses))],
230     y=ffe_train_losses,
231     file_dir='/content/drive/MyDrive/CSE446_hw3/ffe_t.png')
232
233 plot(title='Fixed feature extractor validation loss per epoch',
234     x_label='epoch',
235     y_label='loss',
236     x=[x for x in range(len(ffe_valid_losses))],
237     y=ffe_valid_losses,
238     file_dir='/content/drive/MyDrive/CSE446_hw3/ffe_v.png')
239
240 # Display the test performance of the ffe model
241 loss, accuracy = performance(ffe_model)
242 print(f'test Loss: {loss} Accuracy: {accuracy}')
243
244 # Define the fine tuning model and it's related components.
245 ft_model = torchvision.models.alexnet(pretrained=True)
246
247 # Train every layer.
248 for param in ft_model.parameters():

```

```

248     param.requires_grad = True
249
250 ft_model.classifier[6] = nn.Linear(4096, 10)
251
252 criterion = nn.CrossEntropyLoss()
253 optimizer = optim.SGD(ft_model.parameters(), lr=0.001, momentum=0.9)
254 scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)
255
256 # Train and save the ft model.
257 ft_model, ft_train_losses, ft_valid_losses = train_model(ft_model, criterion, optimizer,
    scheduler, 20)
258
259 PATH = '/content/drive/MyDrive/CSE446_hw3/cifar_ft_model.pth'
260 torch.save(fe_model.state_dict(), PATH)
261
262 plot(title='Fine-tuning training loss per epoch',
263      x_label='epoch',
264      y_label='loss',
265      x=[x for x in range(len(ft_train_losses))],
266      y=ft_train_losses,
267      file_dir='/content/drive/MyDrive/CSE446_hw3/ft_t.png')
268
269 plot(title='Fine-tuning training loss per epoch',
270      x_label='epoch',
271      y_label='loss',
272      x=[x for x in range(len(ft_valid_losses))],
273      y=ft_valid_losses,
274      file_dir='/content/drive/MyDrive/CSE446_hw3/ft_v.png')
275
276 # Display the test performance of the ft model
277 loss, accuracy = performance(ft_model)
278 print(f'test Loss: {loss} Accuracy: {accuracy}')

```