

# Homework #4

CSE 446: Machine Learning

Eric Boris: 1976637

## Conceptual Questions

---

### Problem 1

- True.** There are  $d$  columns in  $X$  and  $k \leq d$ , there are enough columns to project onto a  $k$  dimensional subspace using PCA.
- False.** The columns of  $V$  are equal to the eigenvectors of  $X^T X$ .
- False.** This would create meaningless clusters because it will result in single data point clusters with zero distance from the center of each cluster.
- False.** Consider that  $uIu^T = I$  and  $vIv^T = I$  are distinct singular values of the identity matrix  $I$  where  $u$  and  $v$  are distinct orthogonal matrices.
- False.** Consider that an  $n \times n$  matrix has exactly one eigenvalue.
- True.** Because the autoencoder has non-linear activation functions it can capture more variance than PCA's projection of the data onto a linear subspace.

## Basics of SVD

---

### Problem 2

- Use the SVD of  $X$  to show that  $\hat{w} > \hat{w}_R$ .

$$\begin{aligned}\hat{w} &> \hat{w}_R \\ \min_w \|Xw - y\|_2^2 &> \min_w \|Xw - y\|_2^2 + \lambda \|w\|_2^2 \\ (X^T X)^{-1} X^T y &> (X^T X + \lambda I)^{-1} X^T y \\ (X^T X)^{-1} &> (X^T X + \lambda I)^{-1} \\ (V \Sigma^T U^T U \Sigma V^T)^{-1} &> (V \Sigma^T U^T U \Sigma V^T + \lambda I)^{-1} && \text{Use SVD} \\ VD^{-1}V^T &> V(D + \lambda I)^{-1}V^T && V \text{ is unitary and let } \Sigma^2 = D \text{ s.t. } D \in \mathbb{R}_+^d\end{aligned}$$

Thus, since  $\lambda > 0$ ,  $\hat{w} > \hat{w}_R$ .

- Show that  $UU^T = U^T U = I$ .

$U$  is a square matrix so by the spectral theorem it can be unitarily diagonalized. Let  $U = TIV^T$  such that  $T, I, V \in \mathbb{R}^{n \times n}$  where  $T, V$  have singular values equal to 1 and  $I$  is the identity matrix. Because  $T$  and  $V$  are unitary,

$$TT^T = T^T T = VV^T = V^T V = I$$

Therefore

$$UU^T = TIV^T VIT^T = TIIIT^T = TT^T = I$$

and

$$U^T U = VIT^T TIV^T = IT^T TI = T^T T = I$$

Thus

$$UU^T = U^T U = I$$

c. Show that  $\|Ux\|_2 = \|x\|_2$  for any  $x \in \mathbb{R}^n$ .

Let  $x \in \mathbb{R}^n$ . Then

$$\|Ux\|_2^2 = (Ux)^T (Ux) = x^T U^T U x = x^T x = \|x\|_2^2$$

Taking the square root, it follows that

$$\|Ux\|_2 = \|x\|_2$$

## K-Means Clustering

---

### Problem 3: Answers

- a. See Problem 3: Code
- b. See Figures 1 and 2




Figure 1

- c. See Figure 3

### Problem 3: Code

```
1 # HW4 Problem 3 - K-Means Clustering
2
3 from mnist import MNIST
4 import numpy as np
5 import matplotlib.pyplot as plt
```

Visualization of cluster centers




Figure 2




Figure 3

```

6 import pickle
7
8 def load_dataset(path):
9     ''' Load the mnist handwritten digits dataset. '''
10     mndata = MNIST(path)

```

```

11     X_train, y_train = map(np.array, mndata.load_training())
12     X_test, y_test = map(np.array, mndata.load_testing())
13     X_train = X_train / 255.0
14     X_test = X_test / 255.0
15     return X_train, y_train, X_test, y_test
16
17 def save(output, path):
18     ''' Save the output to the file path. '''
19     with open(path, 'wb') as f:
20         pickle.dump(output, f)
21
22 def plot(*series, title, x_label, y_label, file_path, x_ticks=None):
23     ''' Plot the data. '''
24     plt.title(title)
25     for x, y, label in series:
26         if not x:
27             plt.plot(y, label=label, marker='o')
28         else:
29             plt.plot(x, y, label=label, marker='o')
30     plt.xlabel(x_label)
31     plt.ylabel(y_label)
32     if x_ticks:
33         plt.xticks(x_ticks)
34     plt.legend()
35     plt.savefig(file_path)
36     plt.show()
37
38 def visualize(title, centers, file_path, fig_size=(8, 5), dim=(28, 28), font_size=16):
39     ''' Visualize the cluster centers. '''
40     fig, ax = plt.subplots(2, 5, figsize=fig_size)
41     for j in range(len(centers)):
42         ax.ravel()[j].imshow(centers[j, :].reshape(dim))
43         ax.ravel()[j].axis('off')
44     fig.suptitle(title, fontsize=font_size)
45     plt.savefig(file_path)
46     plt.show()
47
48 class K_Means:
49     def train(self, X, k, conv_distance=0.01, verbose=False):
50         '''
51         Use Lloyds algorithm to compute K-Means clusters over the data X and
52         return the resultant centers and each iteration's measured objective values.
53
54         Arguments:
55             X -- Numpy array of data
56             k -- Integer of clusters to compute
57             conv_distance -- Float minimum distance for determining convergence
58             verbose -- Boolean display progress output if true
59
60         Returns:
61             centers -- k length array of computed center values
62             objectives -- Array of objective values per iteration
63         '''
64         if verbose:
65             print(f'Training with {k} clusters')
66
67         self.centers = X[np.random.choice(np.arange(len(X)), size=k, replace=False)]
68
69         # Only used for plotting results.
70         objectives = []
71
72         converged = False
73         while not converged:
74             # Compute the point distribution
75             point_dist, objective = self._get_point_distribution(X)
76
77             # Move each of the k centers closer to the center of mass.
78             prev_centers = self.centers
79             self.centers = np.array([np.average(X[i], axis=0) for i in point_dist])

```

```

80         # Check whether the centers have converged.
81         max_distance = np.max(np.sum((prev_centers - self.centers) ** 2, axis=1))
82         converged = max_distance < conv_distance
83
84         if verbose:
85             print(max_distance)
86
87         objectives.append(objective)
88
89         return self.centers, objectives
90
91     def _get_point_distribution(self, X):
92         '''
93         Return a distribution of points from X such that for each point i in X
94         i is assigned to the closest of the k centers.
95
96         Arguments:
97             X -- Numpy array of data
98         Returns:
99             point_dist -- Nested list; distribution of points from X onto each of k points
100             objective -- Float; the measured objective for this iteration
101         '''
102         objective = 0
103         k = len(self.centers)
104
105         point_dist = [[] for _ in range(k)]
106
107         # Compute the sum of euclidean distances between two arrays.
108         distance = lambda x, y: np.sum((x - y) ** 2)
109
110         # Find which center j each point i is closest to.
111         for i in range(len(X)):
112             distances = [distance(X[i], self.centers[j]) for j in range(k)]
113             j = np.argmin(distances)
114             point_dist[j].append(i)
115             objective += distances[j]
116
117         return point_dist, objective
118
119     def error(self, X):
120         ''' Compute the error of the trained model on the dataset X. '''
121         objective = 0
122
123         for i in X:
124             min_distance = float('inf')
125             for j in self.centers:
126                 distance = np.linalg.norm((i - j) ** 2)
127                 min_distance = min(distance, min_distance)
128
129             objective += min_distance
130
131         return objective / len(X)
132
133 if __name__ == '__main__':
134     print('Loading data')
135     X_train, y_train, X_test, y_test = load_dataset('../data/python-mnist/data/')
136
137     # Run the algorithm on the MNIST training dataset with k = 10.
138     model = K_Means()
139     centers, objectives = model.train(X_train, k=10, verbose=True)
140
141     # Save the results.
142     save(centers, path='../data/a3_part_b_centers.pickle')
143     save(objectives, path='../data/a3_part_b_objectives.pickle')
144
145     # Plot the objective as a function of the iteration number.
146     plot((None, objectives, 'objective'),
147          title='Objective value per iteration number',
148          xlabel='iteration number',

```

```

149     y_label='objective value',
150     file_path='../figures/a3_plot_b.pdf')
151
152     # Visualize the cluster centers as a 28 x 28 image.
153     visualize(title='Visualization of cluster centers',
154              centers=centers,
155              file_path='../figures/a3_clusters.pdf')
156
157     # For k = {2, 4, 8, 16, 32, 64} run the algorithm on the training dataset to obtain centers.
158     k_vals = [2, 4, 8, 16, 32, 64]
159
160     train_errors = []
161     test_errors = []
162
163     for k in k_vals:
164         model = K_Means()
165         centers, objectives = model.train(X_train, k, verbose=True)
166         train_errors.append(model.error(X_train))
167         test_errors.append(model.error(X_test))
168
169     # Save the results.
170     save(train_errors, path='../data/a3_part_c_train_errors.pickle')
171     save(test_errors, path='../data/a3_part_c_test_errors.pickle')
172
173     # Plot the training and test error as a function of k.
174     plot((k_vals, train_errors, 'train'),
175          (k_vals, test_errors, 'test'),
176          title='Training vs test errors as a function of k',
177          x_label='k',
178          y_label='error',
179          x_ticks=k_vals,
180          file_path='../figures/a3_plot_c.pdf')

```

## PCA

---

### Problem 4: Answers

- a.  $\lambda_1$ : 5.148333441485338  
 $\lambda_2$ : 3.7299894906022617  
 $\lambda_{10}$ : 1.250010757212031  
 $\lambda_{30}$ : 0.36492647617793883  
 $\lambda_{50}$ : 0.16962131566853172  
 $\sum_{i=1}^d \lambda_i$ : 52.83384400094484

- b. Show a general formula for the rank- $k$  PCA approximation of  $x$ .

$$\begin{aligned}
 \Sigma &= \frac{1}{n} (X_{\text{train}} - \mathbf{1}\mu^T)^T (X_{\text{train}} - \mathbf{1}\mu^T) \\
 &= \frac{1}{n} (USV^T)^T (USV^T) && \text{Substitute SVD of } X_{\text{train}} - \mathbf{1}\mu^T \\
 &= \frac{1}{n} VSU^TUS^TV^T \\
 &= \frac{1}{n} VS^2V^T \\
 &= \frac{1}{n} (X_{\text{train}} - \mathbf{1}\mu^T) (X_{\text{train}} - \mathbf{1}\mu^T)^T V \\
 &= \frac{1}{n} VS^2 \\
 &= \frac{1}{n} VD
 \end{aligned}$$

Now consider that since  $S$  is a diagonal matrix,

$$S^2 = \begin{bmatrix} s_1^2 & 0 & 0 & \dots \\ 0 & s_2^2 & \dots & 0 \\ 0 & 0 & \dots & \dots \\ 0 & 0 & \dots & 0 \end{bmatrix}$$

Which implies that

$$\frac{1}{n}D = \begin{bmatrix} \frac{1}{n}d_1 & 0 & 0 & \dots \\ 0 & \frac{1}{n}d_2 & \dots & 0 \\ 0 & 0 & \dots & \dots \\ 0 & 0 & \dots & 0 \end{bmatrix}$$

Therefore, let  $\lambda_i = \frac{1}{n}d_i$  where  $\lambda_i$  is the  $i$ -th eigenvalue of  $V$

$$\begin{aligned} \Sigma V &= \frac{1}{n}VD \\ &= V \begin{bmatrix} \lambda_1 & 0 & 0 & \dots \\ 0 & \lambda_2 & \dots & 0 \\ 0 & 0 & \dots & \dots \\ 0 & 0 & \dots & 0 \end{bmatrix} \\ \Sigma v_i &= \lambda_i v_i \end{aligned}$$

So  $(X_{\text{train}} - \mathbf{1}\mu^T)$  where  $x_i \in \mathbb{R}$  can be projected onto  $k$  dimensions by using the first  $k$  eigenvectors of  $V$  to compute the minimum reconstruction error. Thus, a general formula for the rank- $k$  PCA approximation for  $x$  can be written

$$\bar{x} = V_k V_k^T (x_i - \mu^T) + \mu$$

c. See Figures 4 and 5

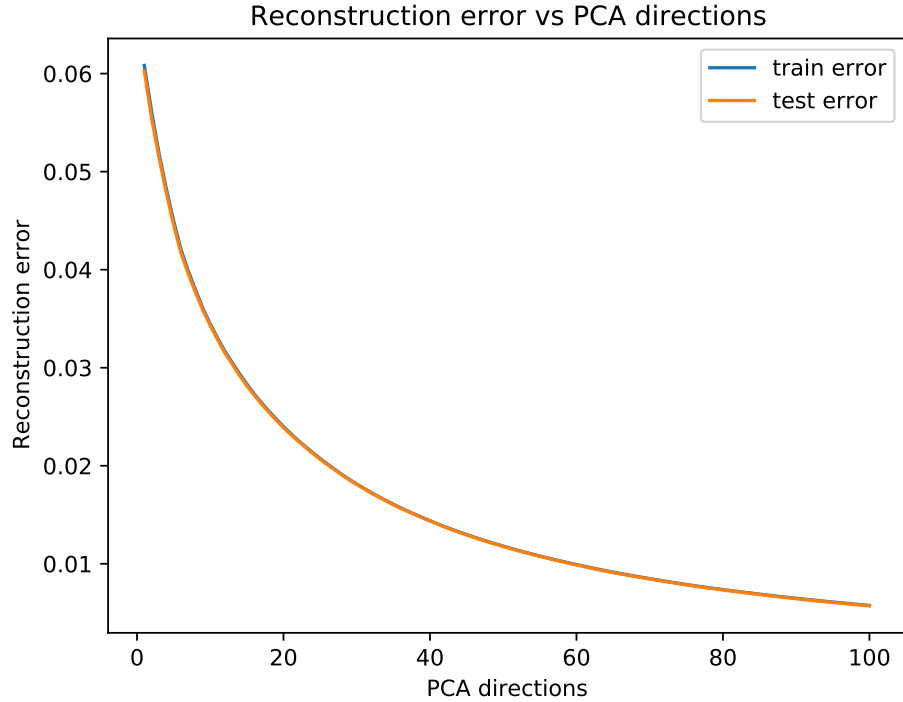


Figure 4




Figure 5

- d. See Figure 6. These represent the first 10 eigenvectors in the single value decomposition of the sample covariance of training examples. They correspond to the most significant entries in the training dataset and so contain overlapping, repeated, and missing digits.





Figure 6

- e. See Figure 7. Increasing  $k$  improves the fidelity of the reconstruction.

#### Problem 4: Code



Figure 7



```
1 # HW4 Problem 4 - PCA
2
3 from mnist import MNIST
```

```

4 import matplotlib.pyplot as plt
5 import numpy as np
6 from tqdm import tqdm
7
8 TQDM_FORM = '{l_bar}{bar:10}{r_bar}{bar:-10b}'
9
10 def load_dataset(path):
11     ''' Load the mnist handwritten digits dataset. '''
12     mndata = MNIST(path)
13     X_train, y_train = map(np.array, mndata.load_training())
14     X_test, y_test = map(np.array, mndata.load_testing())
15     X_train = X_train / 255.0
16     X_test = X_test / 255.0
17     return X_train, y_train, X_test, y_test
18
19 def split(arr, index):
20     ''' Split numpy array into separate sets on the given index. '''
21     return arr[:index, :], arr[index:, :]
22
23 def reconstruct(X, V, mu, k):
24     '''
25     Return the PCA reconstruction from the principle components.
26
27     Arguments:
28         X -- Numpy array of data
29         V -- Numpy array of components
30         mu -- Float of the data mean
31         k -- Integer of top components to use
32
33     Returns:
34         Numpy array of reconstructed data
35     '''
36     # Include only top k components.
37     V = V[:k + 1].T
38     pc_scores = np.matmul(X - mu, V)
39     eigenvectors = np.matmul(pc_scores, V.T)
40     return eigenvectors + mu
41
42 def error(X, V, mu, k):
43     '''
44     Return the mean-squared reconstruction errors of the data over range k.
45
46     Arguments:
47         X -- Numpy array of data
48         V -- Numpy array of components
49         mu -- Float of the data mean
50         k -- Integer of top components to use
51     '''
52     error = []
53     for i in tqdm(range(k), bar_format=TQDM_FORM):
54         X_hat = reconstruct(X, V, mu, i)
55         error.append(np.square(X - X_hat).mean())
56     return error
57
58 def fractional_reconstruction_error(eigenvalues, k):
59     ''' Return a k length array of fractional reconstruction error of the eigenvalues. '''
60     fre = []
61     for i in tqdm(range(k), bar_format=TQDM_FORM):
62         fre.append(1. - np.sum(eigenvalues[:i + 1]) / np.sum(eigenvalues))
63     return fre
64
65 def matching_index(X, digit):
66     ''' Find the first index in the data where the image matches the digit. '''
67     for i in tqdm(range(len(X)), bar_format=TQDM_FORM):
68         if X[i] == digit:
69             break
70     return i
71
72 def plot(*series, title, x_label, y_label, save_path):
73     ''' Plot the data. '''

```

```

73     plt.title(title)
74     for x, y, label in series:
75         plt.plot(x, y, label=label)
76     plt.xlabel(x_label)
77     plt.ylabel(y_label)
78     plt.legend()
79     plt.savefig(save_path)
80     plt.show()
81
82 def visualize(data, iterable, save_path=None, n_rows=2, n_cols=5, dim=(28, 28)):
83     ''' Visualize the data. '''
84     # Minimize the margins
85     plt.gca().set_axis_off()
86     plt.subplots_adjust(top=1, bottom=0, right=1, left=0, hspace=0, wspace=0)
87     plt.margins(0, 0)
88     plt.gca().xaxis.set_major_locator(plt.NullLocator())
89     plt.gca().yaxis.set_major_locator(plt.NullLocator())
90
91     for i, j in enumerate(iterable):
92         plt.subplot(n_rows, n_cols, i + 1)
93         imgplot = plt.imshow(data[j].reshape(dim))
94         plt.axis('off')
95     plt.tight_layout()
96     if save_path:
97         plt.savefig(save_path, bbox_inches='tight', pad_inches=0)
98     plt.show()
99
100 if __name__ == '__main__':
101     print('Load data')
102     X_train, y_train, X_test, y_test = load_dataset('../data/python-mnist/data/')
103     X_train, X_test = split(X_train, index=50000)
104
105     mu = np.mean(X_train, axis=0)
106
107     print('Get eigenvalues')
108     U, S, V = np.linalg.svd(X_train - mu, False)
109     n = X_train.shape[0]
110     eigenvalues = S ** 2 / n
111
112     print('\nPart a')
113     print('What are the eigenvalues 1, 2, 10, 30, and 50?')
114     for i in [0, 1, 9, 29, 49]:
115         print(f'lambda{i + 1}: {eigenvalues[i]}')
116
117     print('What is the sum of eigenvalues?')
118     print(f'Sum of eigenvalues: {np.sum(eigenvalues)}')
119
120     print('\nPart c')
121     print('Compute reconstruction error')
122     train_error = error(X_train, V, mu, k=100)
123     test_error = error(X_test, V, mu, k=100)
124
125     print('Plot reconstruction error')
126     plot((range(1, 101), train_error, 'train error'),
127         (range(1, 101), test_error, 'test error'),
128         title='Reconstruction error vs PCA directions',
129         x_label='PCA directions',
130         y_label='Reconstruction error',
131         save_path='../figures/a4_re.pdf')
132
133     print('Compute fractional reconstruction error')
134     fre = fractional_reconstruction_error(eigenvalues, k=100)
135
136     print('Plot fractional reconstruction error')
137     plot((range(1, 101), fre, 'FRE'),
138         title='Fractional reconstruction error vs PCA directions',
139         x_label='PCA directions',
140         y_label='Fractional reconstruction error',
141         save_path='../figures/a4_fre.pdf')

```

```

142
143     print('\nPart d')
144     print('Display the first 10 eigenvectors')
145     visualize(V, range(0, 10),
146             save_path='../figures/a4_eigenvectors.pdf')
147
148     print('\nPart e')
149     print('Show reconstructions for digits 2, 6, 7 with values k = 5, 15, 40, 100')
150     indices = []
151     for digit in [2, 6, 7]:
152         indices.append(matching_index(y_train, digit))
153
154     # Display the actual digits.
155     visualize(X_train, indices,
156             save_path=f'../figures/a4_actual.pdf',
157             n_rows=1,
158             n_cols=3)
159
160     # Display the reconstructions for different k.
161     for k in [5, 15, 40, 100]:
162         reconstruction = reconstruct(X_train, V, mu, k)
163         visualize(reconstruction, indices,
164                 save_path=f'../figures/a4_recon_{k}.pdf',
165                 n_rows=1,
166                 n_cols=3)
167
168     # Problem 5 Part d.
169     # Re-run PCA with different k to compare results against AutoEncoder.
170
171     # Indices for digits 0-9.
172     digit_indices = [1, 14, 16, 12, 9, 11, 13, 15, 17, 4]
173
174     # Display the actual digits.
175     visualize(X_train, digit_indices,
176             save_path=f'../figures/a5_d_actual.pdf',
177             n_rows=1,
178             n_cols=10)
179
180     # Display the reconstructions for different k.
181     for k in [32, 64, 128]:
182         reconstruction = reconstruct(X_train, V, mu, k)
183         visualize(reconstruction, digit_indices,
184                 save_path=f'../figures/a5_d_recon_{k}.pdf',
185                 n_rows=1,
186                 n_cols=10)

```

## Unsupervised Learning with Autoencoders

---

### Problem 5: Answers

- a. See Figure 8
- b. See Figure 9

Test Errors on Linear and Non-linear networks

c.

	h=32	h=64	h=128
Linear	1096.89	590.39	275.92
Non-linear	958.09	545.69	315.90

- d. See Figure 10. Because of hidden layers and non-linearities, the Autoencoder is better able to capture digit features than the PCA model. Compare and contrast how increasing the number of hidden layers improves the fidelity of the reconstructed digits faster than increasing k in PCA. Further, the Autoencoder

Figure 8: Linear network digit reconstructions

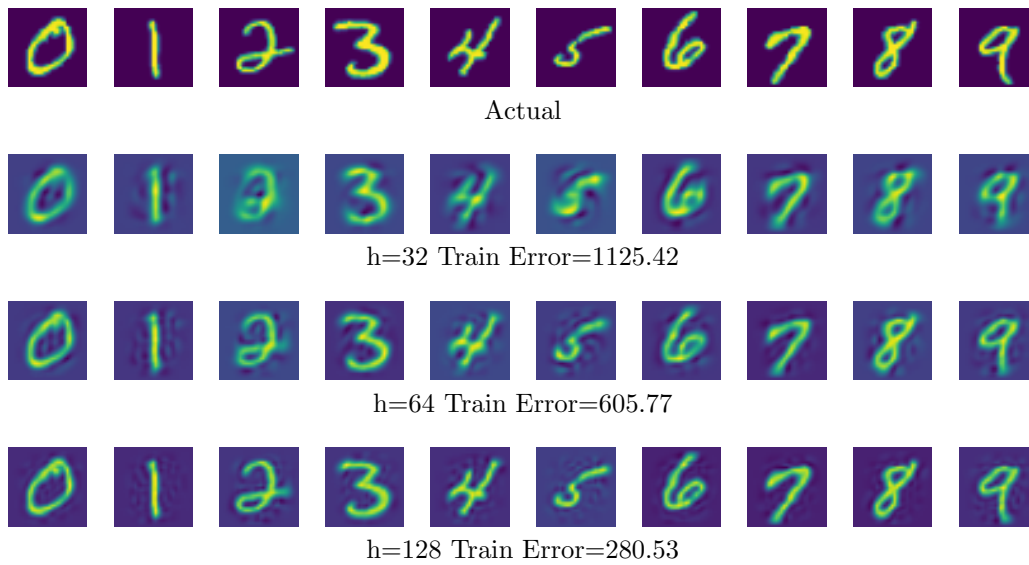



Figure 9: Non-linear network digit reconstructions



uses activation functions for better non-linear boundaries given the non-linear model superior performance over PCA for all values of  $k/h$ .


#### Problem 5: Code

```

1  # HW4 Problem 5 - AutoEncoders
2
3  import torch
4  import torch.nn as nn
5  import torchvision.datasets as datasets
6  from torchvision.datasets import MNIST
7  import matplotlib.pyplot as plt
8  from tqdm import tqdm
9
10 def load_dataset():

```

Figure 10: PCA digit reconstructions



```

11     ''' Load the mnist handwritten digits dataset. '''
12     train = datasets.MNIST(root='.', train=True, download=True, transform=None)
13     test = datasets.MNIST(root='.', train=False, download=True, transform=None)
14     X_train = train.data.view(-1, 784).float()
15     y_train = train.targets
16     X_test = test.data.view(-1, 784).float()
17     y_test = test.targets
18     return X_train, y_train, X_test, y_test
19
20 def visualize(data, iterable, save_path=None, n_rows=1, n_cols=10, dim=(28, 28)):
21     ''' Visualize the data. '''
22     # Minimize the margins
23     plt.gca().set_axis_off()
24     plt.subplots_adjust(top=1, bottom=0, right=1, left=0, hspace=0, wspace=0)
25     plt.margins(0, 0)
26     plt.gca().xaxis.set_major_locator(plt.NullLocator())
27     plt.gca().yaxis.set_major_locator(plt.NullLocator())
28
29     for i, j in enumerate(iterable):
30         plt.subplot(n_rows, n_cols, i + 1)
31         imgplot = plt.imshow(data[j].reshape(dim))
32         plt.axis('off')
33     plt.tight_layout()
34     if save_path:
35         plt.savefig(save_path, bbox_inches='tight', pad_inches=0)
36     plt.show()
37
38 def save(output, path):
39     ''' Save the output to the file path. '''
40     with open(path, 'w') as f:
41         f.write(output)
42
43 class AutoEncoder:
44     def __init__(self, d, h, is_linear=True):
45         self.h = h
46         self.model = self._get_model(d, h, is_linear)
47         self.loss_fn = nn.MSELoss()
48
49     def _get_model(self, d, h, is_linear):
50         ''' Return a linear or non-linear model with the dimensions d and h. '''
51         if is_linear:
52             return torch.nn.Sequential(nn.Linear(d, h), nn.Linear(h, d))

```

```

53         else:
54             return torch.nn.Sequential(nn.Linear(d, h), nn.ReLU(), nn.Linear(h, d), nn.ReLU())
55
56     def train(self, X, n_epochs, learning_rate=1E-3, verbose=False):
57         ''' Train the model and return a list of training losses. '''
58         self.optimizer = torch.optim.Adam(self.model.parameters(), lr=learning_rate)
59         losses = []
60         for i in tqdm(range(n_epochs)):
61             loss = self.get_loss(X)
62             losses.append(loss.item())
63             if verbose:
64                 print(f'iter: {i+1}\tloss: {loss.item()}')
65             self.optimizer.zero_grad()
66             loss.backward()
67             self.optimizer.step()
68         return losses
69
70     def get_loss(self, X):
71         ''' Return the model loss against the given data. '''
72         X_hat = self.model(X)
73         return self.loss_fn(X_hat, X)
74
75 if __name__ == '__main__':
76     # Model parameters
77     d = 784
78     h_vals = [32, 64, 128]
79     n_epochs = 2000
80
81     # Indices for digits 0-9.
82     digit_indices = [1, 14, 16, 12, 9, 11, 13, 15, 17, 4]
83
84     # Determine training device; GPU or CPU.
85     device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
86
87     # Load the datasets.
88     X_train, y_train, X_test, y_test_target = load_dataset()
89
90     # Save the losses.
91     train_losses = ''
92     test_losses = ''
93
94     # Build linear and non-linear models.
95     for is_linear in (True, False):
96         # Build models with d x h dimensions.
97         models = [AutoEncoder(d, h, is_linear) for h in h_vals]
98
99         # Display the actual digits.
100        visualize(X_train, digit_indices, save_path=f'../figures/a5_{is_linear}_actual.pdf')
101
102        for i, m in enumerate(models):
103            print(f'h = {h_vals[i]}')
104
105            # Train the model.
106            losses = m.train(X_train, n_epochs, verbose=False)
107
108            # For parts a and b.
109            # Display each reconstructed digit.
110            with torch.no_grad():
111                visualize(m.model(X_train), digit_indices, save_path=f'../figures/a5_{is_linear}_{h_vals[i]}.pdf')
112
113            # For part c.
114            # Measure the model loss over the test data.
115            test_loss = m.get_loss(X_test)
116
117            # Display the losses.
118            print(f'Train loss: {losses[-1]}')
119            print(f'Test loss: {test_loss.item()}')
120

```

```

121         # Update the loss measures.
122         train_losses += f'{is_linear} {h_vals[i]} {losses[-1]}\n'
123         test_losses += f'{is_linear} {h_vals[i]} {test_loss.item()}\n'
124
125     save(train_losses, path=f'../data/a5_train_losses.txt')
126     save(test_losses, path=f'../data/a5_test_losses.txt')

```

## Text classification on SST-2

---

### Problem 7: Answers

- See Problem 7: Code
- See Problem 7: Code
- See Problem 7: Code

Accuracy and Loss of Neural Networks (8 epochs)

	Training		Validation	
	Accuracy (%)	Loss	Accuracy (%)	Loss
d. RNN	0.8987	0.2719	0.7695	0.5934
GRU	0.9305	0.1866	0.7993	0.6079
LSTM	0.9382	0.1677	0.7259	0.8714

- Output each hidden state. That is, each word produces an output. Then we build a model around predicting the tag of each token from the exposed hidden states.

### Problem 7: Code

```

1  # HW4 Problem 7 - Text Classification
2
3  import torch
4  import torch.nn as nn
5
6  def collate_fn(batch):
7      """
8      Create a batch of data given a list of N sequences and labels.
9      Sequences are stacked into a single tensor of shape (N, max_sequence_length),
10     where max_sequence_length is the maximum length of any sequence in the batch.
11     Sequences shorter than this length should be filled up with 0's.
12     Also returns a tensor of shape (N, 1) containing the label of each sequence.
13
14     :param batch: A list of size N, where each element is a tuple containing a sequence tensor
15     and a single item tensor containing the true label of the sequence.
16
17     :return: A tuple containing two tensors. The first tensor has shape
18     (N, max_sequence_length) and contains all sequences.
19     Sequences shorter than max_sequence_length are padded with 0s at the end.
20     The second tensor has shape (N, 1) and contains all labels.
21     """
22     sentences, labels = zip(*batch)
23     sentences, labels = list(sentences), torch.stack(list(labels))
24
25     max_sequence_length = max(len(s) for s in sentences)
26
27     for i, sentence in enumerate(sentences):
28         length_diff = max_sequence_length - len(sentence)
29         sentences[i] = torch.nn.functional.pad(sentence, [0, length_diff])
30
31     sentences_tensor = torch.stack(sentences)
32     return sentences_tensor, labels
33
34
35 class RNNBinaryClassificationModel(nn.Module):
36     def __init__(self, embedding_matrix, hidden_size=64):

```



```

37     super().__init__()
38     embedding_dim = embedding_matrix.shape[1]
39     self.num_layers = 6
40
41     # Construct embedding layer and initialize with given embedding matrix. Do not modify this code.
42     self.embedding = nn.Embedding(num_embeddings=embedding_matrix.shape[0],
43                                   embedding_dim=embedding_dim,
44                                   padding_idx=0)
45     self.embedding.weight.data = embedding_matrix
46
47     # Construct 3 different types of RNN for comparison.
48     self.RNN = nn.RNN(input_size=embedding_dim,
49                       hidden_size=hidden_size,
50                       num_layers=self.num_layers,
51                       batch_first=True)
52
53     self.GRU = nn.GRU(input_size=embedding_dim,
54                       hidden_size=hidden_size,
55                       num_layers=self.num_layers,
56                       batch_first=True)
57
58     self.LSTM = nn.LSTM(input_size=embedding_dim,
59                        hidden_size=hidden_size,
60                        num_layers=self.num_layers,
61                        batch_first=True,
62                        bidirectional=True)
63
64     self.linear = nn.Linear(in_features=hidden_size, out_features=1)
65     self.sigmoid = nn.Sigmoid()
66
67     def forward(self, inputs):
68         """
69         Takes in a batch of data of shape (N, max_sequence_length).
70         Returns a tensor of shape (N, 1), where each
71         element corresponds to the prediction for the corresponding sequence.
72         :param inputs: Tensor of shape (N, max_sequence_length) containing N
73         sequences to make predictions for.
74         :return: Tensor of predictions for each sequence of shape (N, 1).
75         """
76         # Un-comment for training other models.
77         #return self.sigmoid(self.linear(self.RNN(self.embedding(inputs))[1][-1].squeeze(0)))
78         #return self.sigmoid(self.linear(self.GRU(self.embedding(inputs))[1][-1].squeeze(0)))
79         return self.sigmoid(self.linear(self.LSTM(self.embedding(inputs))[1][0][-1].squeeze(0)))
80
81     def loss(self, logits, targets):
82         """
83         Computes the binary cross-entropy loss.
84         :param logits: Raw predictions from the model of shape (N, 1)
85         :param targets: True labels of shape (N, 1)
86         :return: Binary cross entropy loss between logits and targets as a scalar tensor.
87         """
88         return nn.BCELoss()(logits, targets.float())
89
90     def accuracy(self, logits, targets):
91         """
92         Computes the accuracy, i.e number of correct predictions / N.
93         :param logits: Raw predictions from the model of shape (N, 1)
94         :param targets: True labels of shape (N, 1)
95         :return: Accuracy as a scalar tensor.
96         """
97         correct = 0
98
99         for i in range(len(logits)):
100             prediction = torch.round(logits[i])
101             correct += (prediction == targets[i]).sum().item()
102
103         accuracy = correct / len(logits)
104         return torch.tensor(accuracy)
105

```

```
106 # Training parameters
107 TRAINING_BATCH_SIZE = 32
108 NUM_EPOCHS = 8
109 LEARNING_RATE = 0.0001
110
111 # Batch size for validation, this only affects performance.
112 VAL_BATCH_SIZE = 128
```